

Problem 7

```
In [154]: import numpy as np
import urllib
import scipy.optimize
from random import shuffle
import string
from collections import defaultdict
from nltk import bigrams
from nltk.stem.porter import *
from sklearn import linear_model
from sklearn.metrics import mean_squared_error
from sklearn.metrics.pairwise import cosine_similarity
from math import log
```

```
In [155]: def parseData(fname):
    for l in urllib.request.urlopen(fname):
        yield eval(l)
```

```
In [156]: print("Reading data...")
data = list(parseData("file:beer_50000.json"))
print("done")
```

Reading data...
done

```
In [157]: shuffle(data)
```

```
In [158]: data_train = data[:5000]
data_valid = data[5000:10000]
data_test = data[10000:15000]
```

1. unigrams, removing punctuation, tfidf

```
In [159]: punctuation = set(string.punctuation)
reviewText = [''.join([c for c in datum['review/text'].lower() if not c in
```

```
In [160]: unigramCount = defaultdict(int)
for d in data:
    r = ''.join([c for c in d['review/text'].lower() if not c in punctuati
    for w in r.split():
        unigramCount[w] += 1
```

```
In [161]: countsUnigram = [(unigramCount[w], w) for w in unigramCount]
unigrams = [x[1] for x in countsUnigram[:1000]]
unigramId = dict(zip(unigrams, range(len(unigrams))))
unigramSet = set(unigrams)
```

```
In [162]: countsUnigram.sort()
countsUnigram.reverse()
```

```
In [163]: def feature_000(datum):
    feat = [0] * len(unigrams)
    unigramCount_rt = defaultdict(int)
    r = ''.join([c for c in datum['review/text'].lower() if not c in punct
    for w in r.split():
        unigramCount_rt[w] += 1
    for unigram in unigramId:
        if unigram in unigramCount_rt:
            feat[unigramId[unigram]] = unigramCount_rt[unigram] * log(150
    feat.append(1) #offset
    return feat
```

```
In [164]: X_train = [feature_000(d) for d in data_train]
y_train = [d['review/overall'] for d in data_train]
X_valid = [feature_000(d) for d in data_valid]
y_valid = [d['review/overall'] for d in data_valid]
X_test = [feature_000(d) for d in data_test]
y_test = [d['review/overall'] for d in data_test]
```

```
In [165]: def compare_unigrams(reg):
    clf = linear_model.Ridge(reg, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_
    predictions = clf.predict(X_valid)
    predictions = np.matrix(predictions)
    y_validation = np.matrix(y_valid)
    diff = y_validation - predictions
    MSE = diff * diff.T / 5000
    print("reg for unigrams= ", r, "MSE = ", MSE)
```

```
In [166]: reg = [0.01, 0.1, 1.0, 10, 100]
for r in reg:
    compare_unigrams(r)

reg for unigrams= 0.01 MSE = [[0.40196911]]
reg for unigrams= 0.1 MSE = [[0.39984351]]
reg for unigrams= 1.0 MSE = [[0.39184914]]
reg for unigrams= 10 MSE = [[0.37368687]]
reg for unigrams= 100 MSE = [[0.38978366]]
```

```
In [167]: clf = linear_model.Ridge(10, fit_intercept=False)
clf.fit(X_train, y_train)
theta = clf.coef_
predictions = clf.predict(X_test)
predictions = np.matrix(predictions)
y_testing = np.matrix(y_test)
diff = y_testing - predictions
MSE = diff * diff.T / 5000
print("optimal lambda = 10, testing MSE = " + str(MSE))

optimal lambda = 10, testing MSE = [[0.38394586]]
```

2. unigrams, removing punctuation, word counts

```
In [168]: punctuation = set(string.punctuation)
reviewText = [''.join([c for c in datum['review/text'].lower() if not c in
```

```
In [169]: unigramCount = defaultdict(int)
for d in data:
    r = ''.join([c for c in d['review/text'].lower() if not c in punctuati
    for w in r.split():
        unigramCount[w] += 1
```

```
In [170]: countsUnigram = [(unigramCount[w], w) for w in unigramCount]
unigrams = [x[1] for x in countsUnigram[:1000]]
unigramId = dict(zip(unigrams, range(len(unigrams))))
unigramSet = set(unigrams)
```

```
In [171]: countsUnigram.sort()
countsUnigram.reverse()
```

```
In [172]: def feature_001(datum):
    feat = [0] * len(unigrams)
    unigramCount_rt = defaultdict(int)
    r = ''.join([c for c in datum['review/text'].lower() if not c in punct
    for w in r.split():
        if w in unigrams:
            feat[unigramId[w]] += 1
    feat.append(1) #offset
    return feat
```

```
In [173]: X_train = [feature_001(d) for d in data_train]
y_train = [d['review/overall'] for d in data_train]
X_valid = [feature_001(d) for d in data_valid]
y_valid = [d['review/overall'] for d in data_valid]
X_test = [feature_001(d) for d in data_test]
y_test = [d['review/overall'] for d in data_test]
```

```
In [174]: def compare_unigrams(reg):
    clf = linear_model.Ridge(reg, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_
    predictions = clf.predict(X_valid)
    predictions = np.matrix(predictions)
    y_validation = np.matrix(y_valid)
    diff = y_validation - predictions
    MSE = diff * diff.T / 5000
    print("reg for unigrams= ", r, "MSE = ", MSE)
```

```
In [175]: reg = [0.01, 0.1, 1.0, 10, 100]
          for r in reg:
              compare_unigrams(r)

reg for unigrams= 0.01 MSE = [[0.40064272]]
reg for unigrams= 0.1 MSE = [[0.3948727]]
reg for unigrams= 1.0 MSE = [[0.38137263]]
reg for unigrams= 10 MSE = [[0.36343046]]
reg for unigrams= 100 MSE = [[0.39376672]]
```

```
In [176]: clf = linear_model.Ridge(10, fit_intercept=False)
          clf.fit(X_train, y_train)
          theta = clf.coef_
          predictions = clf.predict(X_test)
          predictions = np.matrix(predictions)
          y_testing = np.matrix(y_test)
          diff = y_testing - predictions
          MSE = diff * diff.T / 5000
          print("optimal lambda = 10, testing MSE = " + str(MSE))

optimal lambda = 10, testing MSE = [[0.37596804]]
```

3. unigrams, preserving punctuation, tfidf

```
In [177]: punctuation = set(string.punctuation)
          reviewText = [''.join([c for c in datum['review/text'].lower() if not c in
```

```
In [178]: unigramCount = defaultdict(int)
          for d in data:
              unigramList = []
              for c in d['review/text'].lower():
                  if c not in punctuation:
                      unigramList.append(c)
                  else:
                      unigramList.append(" ")
                      unigramList.append(c)
              r = ''.join(unigramList)
              for w in r.split():
                  unigramCount[w] += 1
```

```
In [179]: countsUnigram = [(unigramCount[w], w) for w in unigramCount]
          unigrams = [x[1] for x in countsUnigram[:1000]]
          unigramId = dict(zip(unigrams, range(len(unigrams))))
          unigramSet = set(unigrams)
```

```
In [180]: countsUnigram.sort()
          countsUnigram.reverse()
```

```
In [181]: def feature_010(datum):
    feat = [0] * len(unigrams)
    unigramCount_datum = defaultdict(int)

    unigramList = []
    for c in datum['review/text'].lower():
        if c not in punctuation:
            unigramList.append(c)
        else:
            unigramList.append(" ")
            unigramList.append(c)
    r = ''.join(unigramList)

    for w in r.split():
        unigramCount_datum[w] += 1

    for unigram in unigramId:
        if unigram in unigramCount_datum:
            feat[unigramId[unigram]] = unigramCount_datum[unigram] * log(1

    feat.append(1) #offset
    return feat
```

```
In [182]: X_train = [feature_010(d) for d in data_train]
    y_train = [d['review/overall'] for d in data_train]
    X_valid = [feature_010(d) for d in data_valid]
    y_valid = [d['review/overall'] for d in data_valid]
    X_test = [feature_010(d) for d in data_test]
    y_test = [d['review/overall'] for d in data_test]
```

```
In [183]: def compare_unigrams(reg):
    clf = linear_model.Ridge(reg, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_
    predictions = clf.predict(X_valid)
    predictions = np.matrix(predictions)
    y_validation = np.matrix(y_valid)
    diff = y_validation - predictions
    MSE = diff * diff.T / 5000
    print("reg for unigrams= ", r, "MSE = ", MSE)
```

```
In [184]: reg = [0.01, 0.1, 1.0, 10, 100]
    for r in reg:
        compare_unigrams(r)

reg for unigrams= 0.01 MSE = [[0.40508638]]
reg for unigrams= 0.1 MSE = [[0.40344859]]
reg for unigrams= 1.0 MSE = [[0.39609762]]
reg for unigrams= 10 MSE = [[0.37728858]]
reg for unigrams= 100 MSE = [[0.39379484]]
```

```
In [185]: clf = linear_model.Ridge(10, fit_intercept=False)
clf.fit(X_train, y_train)
theta = clf.coef_
predictions = clf.predict(X_test)
predictions = np.matrix(predictions)
y_testing = np.matrix(y_test)
diff = y_testing - predictions
MSE = diff * diff.T / 5000
print("optimal lambda = 10, testing MSE = " + str(MSE))
```

optimal lambda = 10, testing MSE = [[0.38785107]]

4. unigrams, preserving punctuation, word counts

```
In [186]: punctuation = set(string.punctuation)
reviewText = ''.join([c for c in datum['review/text'].lower() if not c in
```

```
In [187]: unigramCount = defaultdict(int)
for d in data:
    unigramList = []
    for c in d['review/text'].lower():
        if c not in punctuation:
            unigramList.append(c)
        else:
            unigramList.append(" ")
            unigramList.append(c)
    r = ''.join(unigramList)
    for w in r.split():
        unigramCount[w] += 1
```

```
In [188]: countsUnigram = [(unigramCount[w], w) for w in unigramCount]
unigrams = [x[1] for x in countsUnigram[:1000]]
unigramId = dict(zip(unigrams, range(len(unigrams))))
unigramSet = set(unigrams)
```

```
In [189]: countsUnigram.sort()
countsUnigram.reverse()
```

```
In [190]: def feature_011(datum):
    feat = [0] * len(unigrams)
    #unigramCount_rt = defaultdict(int)
    r = ''.join([c for c in datum['review/text'].lower() if not c in punct
    for w in r.split():
        if w in unigrams:
            feat[unigramId[w]] += 1
    feat.append(1) #offset
    return feat
```

```
In [191]: X_train = [feature_010(d) for d in data_train]
y_train = [d['review/overall'] for d in data_train]
X_valid = [feature_010(d) for d in data_valid]
y_valid = [d['review/overall'] for d in data_valid]
X_test = [feature_010(d) for d in data_test]
y_test = [d['review/overall'] for d in data_test]
```

```
In [192]: def compare_unigrams(reg):
    clf = linear_model.Ridge(reg, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_
    predictions = clf.predict(X_valid)
    predictions = np.matrix(predictions)
    y_validation = np.matrix(y_valid)
    diff = y_validation - predictions
    MSE = diff * diff.T / 5000
    print("reg for unigrams= ", r, "MSE = ", MSE)
```

```
In [193]: reg = [0.01, 0.1, 1.0, 10, 100]
for r in reg:
    compare_unigrams(r)

reg for unigrams= 0.01 MSE = [[0.40508638]]
reg for unigrams= 0.1 MSE = [[0.40344859]]
reg for unigrams= 1.0 MSE = [[0.39609762]]
reg for unigrams= 10 MSE = [[0.37728858]]
reg for unigrams= 100 MSE = [[0.39379484]]
```

```
In [194]: clf = linear_model.Ridge(10, fit_intercept=False)
clf.fit(X_train, y_train)
theta = clf.coef_
predictions = clf.predict(X_test)
predictions = np.matrix(predictions)
y_testing = np.matrix(y_test)
diff = y_testing - predictions
MSE = diff * diff.T / 5000
print("optimal lambda = 10, testing MSE = " + str(MSE))
```

```
optimal lambda = 10, testing MSE = [[0.38785107]]
```

5. bigrams, removing punctuation, tfidf

```
In [195]: bigramCount = defaultdict(int)
punctuation = set(string.punctuation)
#stemmer = PorterStemmer()
for d in data:
    r = ''.join([c for c in d['review/text'].lower() if not c in punctuation])
    words = r.split()
    for i in range(0, len(words)-1):
        bigram = words[i] + " " + words[i+1]
        bigramCount[bigram] += 1
```

```
In [196]: countsBigram = [(bigramCount[d], d) for d in bigramCount.keys()]
countsBigram.sort()
countsBigram.reverse()
```

```
In [197]: bigrams = [c[1] for c in countsBigram[:1000]]
bigramId = dict(zip(bigrams, range(len(bigrams))))
bigramSet = set(bigrams)
```

```
In [198]: def feature_100(datum):
    feat = [0] * len(bigrams)
    bigramCount_datum = defaultdict(int)
    r = ''.join([c for c in datum['review/text'].lower() if not c in punct
    words = r.split()
    for i in range(len(words) - 1):
        bigram = words[i] + " " + words[i + 1]
        bigramCount_datum[bigram] += 1

    for bigram in bigramId:
        if bigram in bigramCount_datum:
            feat[bigramId[bigram]] = bigramCount_datum[bigram] * log(15000

    feat.append(1) #offset
    return feat
```

```
In [199]: X_train = [feature_100(d) for d in data_train]
y_train = [d['review/overall'] for d in data_train]
X_valid = [feature_100(d) for d in data_valid]
y_valid = [d['review/overall'] for d in data_valid]
X_test = [feature_100(d) for d in data_test]
y_test = [d['review/overall'] for d in data_test]
```

```
In [200]: def compare_bigrams(reg):
    clf = linear_model.Ridge(reg, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_
    predictions = clf.predict(X_valid)
    predictions = np.matrix(predictions)
    y_validation = np.matrix(y_valid)
    diff = y_validation - predictions
    MSE = diff * diff.T / 5000
    print("reg for bigrams= ", r, "MSE = ", MSE)
```

```
In [201]: reg = [0.01, 0.1, 1.0, 10, 100]
for r in reg:
    compare_bigrams(r)

reg for bigrams= 0.01 MSE = [[0.44638918]]
reg for bigrams= 0.1 MSE = [[0.44612483]]
reg for bigrams= 1.0 MSE = [[0.44349019]]
reg for bigrams= 10 MSE = [[0.42654711]]
reg for bigrams= 100 MSE = [[0.4132646]]
```



```
In [202]: clf = linear_model.Ridge(100, fit_intercept=False)
clf.fit(X_train, y_train)
theta = clf.coef_
predictions = clf.predict(X_test)
predictions = np.matrix(predictions)
y_testing = np.matrix(y_test)
diff = y_testing - predictions
MSE = diff * diff.T / 5000
print("optimal lambda = 100, testing MSE = " + str(MSE))
```

```
optimal lambda = 100, testing MSE = [[0.42533635]]
```

6. bigrams, removing punctuation, word counts

```
In [203]: bigramCount = defaultdict(int)
punctuation = set(string.punctuation)
#stemmer = PorterStemmer()
for d in data:
    r = ''.join([c for c in d['review/text'].lower() if not c in punctuation])
    words = r.split()
    for i in range(0, len(words)-1):
        bigram = words[i] + " " + words[i+1]
        bigramCount[bigram] += 1
```

```
In [204]: countsBigram = [(bigramCount[d], d) for d in bigramCount.keys()]
countsBigram.sort()
countsBigram.reverse()
```

```
In [205]: bigrams = [c[1] for c in countsBigram[:1000]]
bigramId = dict(zip(bigrams, range(len(bigrams))))
bigramSet = set(bigrams)
```

```
In [206]: def feature_101(datum):
    feat = [0]*len(bigrams)
    r = ''.join([c for c in datum['review/text'].lower() if not c in punctuation])
    words = r.split()
    for i in range(len(words)-1):
        bigram = words[i] + " " + words[i + 1]
        if bigram in bigrams:
            feat[bigramId[bigram]] += 1
    feat.append(1) #offset
    return feat
```

```
In [207]: X_train = [feature_101(d) for d in data_train]
y_train = [d['review/overall'] for d in data_train]
X_valid = [feature_101(d) for d in data_valid]
y_valid = [d['review/overall'] for d in data_valid]
X_test = [feature_101(d) for d in data_test]
y_test = [d['review/overall'] for d in data_test]
```

```
In [208]: def compare_bigrams(reg):
    clf = linear_model.Ridge(reg, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_
    predictions = clf.predict(X_valid)
    predictions = np.matrix(predictions)
    y_validation = np.matrix(y_valid)
    diff = y_validation - predictions
    MSE = diff * diff.T / 5000
    print("reg for bigrams= ", r, "MSE = ", MSE)
```

```
In [209]: reg = [0.01, 0.1, 1.0, 10, 100]
    for r in reg:
        compare_bigrams(r)

reg for bigrams= 0.01 MSE = [[0.44597822]]
reg for bigrams= 0.1 MSE = [[0.44569662]]
reg for bigrams= 1.0 MSE = [[0.44299749]]
reg for bigrams= 10 MSE = [[0.42425457]]
reg for bigrams= 100 MSE = [[0.41755767]]
```

```
In [210]: clf = linear_model.Ridge(100, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_
    predictions = clf.predict(X_test)
    predictions = np.matrix(predictions)
    y_testing = np.matrix(y_test)
    diff = y_testing - predictions
    MSE = diff * diff.T / 5000
    print("optimal lambda = 100, testing MSE = " + str(MSE))
```

```
optimal lambda = 100, testing MSE = [[0.43058102]]
```

7. bigrams, preserving punctuation, tfidf

```
In [211]: punctuation = set(string.punctuation)
    reviewText = [''.join([c for c in datum['review/text'].lower() if not c in
```

```
In [212]: bigramCount = defaultdict(int)
    for d in data:
        bigramList = []
        for c in d['review/text'].lower():
            if c not in punctuation:
                bigramList.append(c)
            else:
                bigramList.append(" ")
                bigramList.append(c)
        r = ''.join(bigramList)
        words = r.split()
        for i in range(len(words) - 1):
            bigram = words[i] + " " + words[i + 1]
            bigramCount[bigram] += 1
```

```
In [213]: countsBigram = [(bigramCount[d], d) for d in bigramCount.keys()]
countsBigram.sort()
countsBigram.reverse()
```

```
In [214]: bigrams = [c[1] for c in countsBigram[:1000]]
bigramId = dict(zip(bigrams, range(len(bigrams))))
bigramSet = set(bigrams)
```

```
In [215]: def feature_110(datum):
    feat = [0] * len(bigrams)
    bigramCount_datum = defaultdict(int)

    bigramList = []
    for c in datum['review/text'].lower():
        if c not in punctuation:
            bigramList.append(c)
        else:
            bigramList.append(" ")
            bigramList.append(c)
    r = ''.join(bigramList)
    words = r.split()
    for i in range(len(words) - 1):
        bigram = words[i] + " " + words[i + 1]
        bigramCount_datum[bigram] += 1

    for bigram in bigramId:
        if bigram in bigramCount_datum:
            feat[bigramId[bigram]] = bigramCount_datum[bigram] * log(15000

    feat.append(1) #offset
    return feat
```

```
In [216]: X_train = [feature_110(d) for d in data_train]
y_train = [d['review/overall'] for d in data_train]
X_valid = [feature_110(d) for d in data_valid]
y_valid = [d['review/overall'] for d in data_valid]
X_test = [feature_110(d) for d in data_test]
y_test = [d['review/overall'] for d in data_test]
```

```
In [217]: def compare_bigrams(reg):
    clf = linear_model.Ridge(reg, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_
    predictions = clf.predict(X_valid)
    predictions = np.matrix(predictions)
    y_validation = np.matrix(y_valid)
    diff = y_validation - predictions
    MSE = diff * diff.T / 5000
    print("reg for bigrams= ", r, "MSE = ", MSE)
```

```
In [218]: reg = [0.01, 0.1, 1.0, 10, 100]
          for r in reg:
              compare_bigrams(r)

reg for bigrams= 0.01 MSE = [[0.44922545]]
reg for bigrams= 0.1 MSE = [[0.44880617]]
reg for bigrams= 1.0 MSE = [[0.44514748]]
reg for bigrams= 10 MSE = [[0.42573408]]
reg for bigrams= 100 MSE = [[0.41411603]]
```

```
In [219]: clf = linear_model.Ridge(100, fit_intercept=False)
          clf.fit(X_train, y_train)
          theta = clf.coef_
          predictions = clf.predict(X_test)
          predictions = np.matrix(predictions)
          y_testing = np.matrix(y_test)
          diff = y_testing - predictions
          MSE = diff * diff.T / 5000
          print("optimal lambda = 100, testing MSE = " + str(MSE))

optimal lambda = 100, testing MSE = [[0.4299674]]
```

8. bigrams, preserving punctuation, word counts

```
In [220]: punctuation = set(string.punctuation)
          reviewText = ''.join([c for c in datum['review/text'].lower() if not c in
```

```
In [221]: bigramCount = defaultdict(int)
          for d in data:
              bigramList = []
              for c in d['review/text'].lower():
                  if c not in punctuation:
                      bigramList.append(c)
                  else:
                      bigramList.append(" ")
                      bigramList.append(c)
              r = ''.join(bigramList)
              words = r.split()
              for i in range(len(words) - 1):
                  bigram = words[i] + " " + words[i + 1]
                  bigramCount[bigram] += 1
```

```
In [222]: countsBigram = [(bigramCount[d], d) for d in bigramCount.keys()]
          countsBigram.sort()
          countsBigram.reverse()
```

```
In [223]: bigrams = [c[1] for c in countsBigram[:1000]]
          bigramId = dict(zip(bigrams, range(len(bigrams))))
          bigramSet = set(bigrams)
```

```
In [224]: def feature_111(datum):
    feat = [0] * len(bigrams)
    # bigramCount_datum = defaultdict(int)

    bigramList = []
    for c in datum['review/text'].lower():
        if c not in punctuation:
            bigramList.append(c)
        else:
            bigramList.append(" ")
            bigramList.append(c)
    r = ''.join(bigramList)
    words = r.split()
    for i in range(len(words) - 1):
        bigram = words[i] + " " + words[i + 1]
        if bigram in bigrams:
            feat[bigramId[bigram]] += 1

    feat.append(1) #offset
    return feat
```

```
In [225]: X_train = [feature_111(d) for d in data_train]
y_train = [d['review/overall'] for d in data_train]
X_valid = [feature_111(d) for d in data_valid]
y_valid = [d['review/overall'] for d in data_valid]
X_test = [feature_111(d) for d in data_test]
y_test = [d['review/overall'] for d in data_test]
```

```
In [226]: def compare_bigrams(reg):
    clf = linear_model.Ridge(reg, fit_intercept=False)
    clf.fit(X_train, y_train)
    theta = clf.coef_
    predictions = clf.predict(X_valid)
    predictions = np.matrix(predictions)
    y_validation = np.matrix(y_valid)
    diff = y_validation - predictions
    MSE = diff * diff.T / 5000
    print("reg for bigrams= ", r, "MSE = ", MSE)
```

```
In [227]: reg = [0.01, 0.1, 1.0, 10, 100]
for r in reg:
    compare_bigrams(r)

reg for bigrams= 0.01 MSE = [[0.44906436]]
reg for bigrams= 0.1 MSE = [[0.4487985]]
reg for bigrams= 1.0 MSE = [[0.44627841]]
reg for bigrams= 10 MSE = [[0.42936252]]
reg for bigrams= 100 MSE = [[0.42557995]]
```

```
In [228]: clf = linear_model.Ridge(100, fit_intercept=False)
          clf.fit(X_train, y_train)
          theta = clf.coef_
          predictions = clf.predict(X_test)
          predictions = np.matrix(predictions)
          y_testing = np.matrix(y_test)
          diff = y_testing - predictions
          MSE = diff * diff.T / 5000
          print("optimal lambda = 100, testing MSE = " + str(MSE))
```

```
optimal lambda = 100, testing MSE = [[0.44222335]]
```